

Tutorial 10


Autoencoders

Chen-Change Loy

吕健勤

<https://www.mmlab-ntu.com/person/ccloy/>

<https://twitter.com/ccloy>



Question 1a

Given five binary patterns:

$$\mathbf{x}_1 = \begin{pmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}, \mathbf{x}_2 = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{pmatrix}, \mathbf{x}_3 = \begin{pmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{pmatrix}, \mathbf{x}_4 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{x}_5 = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix}$$

Design an autoencoder with four hidden neurons to reconstruct the patterns, using gradient descent learning with a learning parameter $\alpha = 0.1$.

Find the weights, biases, hidden-layer activations and reconstructions of the input patterns at convergence.

Repeat the above by introducing a sparsity constraint with a penalty parameter $\beta = 0.5$ and sparsity parameter $\rho = 0.1$.

[Recap] Training autoencoders

If the inputs are interpreted as **bit vectors** or **vectors of bit probabilities**, **cross-entropy of the reconstruction** can be used:

$$J_{cross-entropy} = - \sum_{p=1}^P (x_p \log y_p + (1 - x_p) \log(1 - y_p))$$

Learning are done by using gradient descent:

$$\mathbf{W} \leftarrow \mathbf{W} - \alpha \nabla_{\mathbf{W}} J$$

$$\mathbf{b} \leftarrow \mathbf{b} - \alpha \nabla_{\mathbf{b}} J$$

$$\mathbf{c} \leftarrow \mathbf{c} - \alpha \nabla_{\mathbf{c}} J$$

Question 1

```
[17] # Display weights and biases
print(f'W:\n {autoencoder.W.data}\n')
print(f'W_prime:\n {autoencoder.W_prime.data}\n')
print(f'b:\n {autoencoder.b.data}\n')
print(f'b_prime:\n {autoencoder.b_prime.data}\n')
```

```
W:
tensor([[ 1.3404,  1.4979,  1.1027, -1.6822],
        [ 0.7539, -0.4255,  0.8947, -2.8855],
        [-0.6857, -2.6170,  2.0785, -0.5724],
        [ 2.4732, -0.0064, -2.3399,  1.0806],
        [-3.8564,  0.2799,  1.9536,  1.9443],
        [ 0.0641,  0.3798,  0.0269, -0.2577],
        [ 1.1192, -2.2494, -0.7006,  2.4961],
        [ 0.4663,  1.2535, -2.7378,  1.2046],
        [-0.8174,  3.8036, -1.1091,  0.7902]])

W_prime:
tensor([[ 4.5253,  0.8765, -2.3588,  4.4418, -6.2649, -1.2592,  1.4582,  0.7749,
         -2.4911],
        [ 5.3236, -2.2052, -6.1462,  0.9732,  0.2411, -1.1617, -4.2555,  3.2673,
         6.7966],
        [-1.0380,  2.4406,  4.2446, -5.9221,  2.3335, -1.6839, -3.9068, -6.1726,
        -2.6724],
        [-2.4691, -6.0329, -2.7513,  1.6180,  3.1008, -1.5105,  5.9788,  0.7012,
         1.2132]])

b:
tensor([-0.4316, -0.8084,  1.3958,  0.3283])

b_prime:
tensor([ 0.9000, -0.0059,  1.2902, -0.8071, -1.4364, -2.9330,  0.3787, -2.1170,
        -1.8984])
```

```
[18] # Evaluate result of reconstruction
with torch.no_grad():
    h, y, o = autoencoder(X)
print(f'Input:\n {X}\n')
print(f'Hidden activation:\n {h}\n')
print(f'Output:\n {y}\n')
print(f'Output_binary:\n {o}\n')
```



```
Input:
tensor([[1., 1., 1., 0., 0., 0., 0., 0., 0.],
        [1., 0., 0., 1., 0., 0., 1., 0., 0.],
        [0., 0., 1., 0., 1., 0., 1., 0., 0.],
        [1., 0., 0., 0., 1., 0., 0., 0., 1.],
        [1., 0., 0., 1., 0., 0., 1., 1., 1.]])
```

```
Hidden activation:
tensor([[0.7265, 0.0868, 0.9958, 0.0081],
        [0.9890, 0.1727, 0.3677, 0.9023],
        [0.0207, 0.0045, 0.9912, 0.9852],
        [0.0226, 0.9916, 0.9659, 0.7991],
        [0.9845, 0.9704, 0.0123, 0.9855]])
```

```
Output:
tensor([[9.7331e-01, 9.4381e-01, 9.6258e-01, 3.2928e-02, 2.6134e-02, 3.5483e-03,
        5.8771e-02, 6.0390e-04, 3.1119e-03],
        [9.7554e-01, 1.6859e-02, 4.6263e-02, 9.5418e-01, 1.9171e-02, 1.7244e-03,
        9.9359e-01, 8.1404e-02, 4.4102e-02],
        [7.9918e-02, 2.8694e-02, 9.3763e-01, 6.7819e-03, 9.7818e-01, 2.1902e-03,
        9.1739e-01, 5.4520e-04, 3.3148e-02],
        [9.6463e-01, 9.6007e-03, 4.9424e-02, 1.5239e-02, 9.6747e-01, 9.6071e-04,
        5.7109e-02, 1.3909e-02, 9.5981e-01],
        [9.9969e-01, 7.4719e-04, 6.4053e-05, 9.9761e-01, 1.3577e-02, 1.1024e-03,
        9.7149e-01, 9.1922e-01, 9.6794e-01]])
```

```
Output_binary:
tensor([[1., 1., 1., 0., 0., 0., 0., 0., 0.],
        [1., 0., 0., 1., 0., 0., 1., 0., 0.],
        [0., 0., 1., 0., 1., 0., 1., 0., 0.],
        [1., 0., 0., 0., 1., 0., 0., 0., 1.],
        [1., 0., 0., 1., 0., 0., 1., 1., 1.]])
```

Question 1



Question 1

Given five binary patterns:

$$\mathbf{x}_1 = \begin{pmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{pmatrix}, \mathbf{x}_2 = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 0 & 0 \end{pmatrix}, \mathbf{x}_3 = \begin{pmatrix} 0 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \end{pmatrix}, \mathbf{x}_4 = \begin{pmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{pmatrix}, \mathbf{x}_5 = \begin{pmatrix} 1 & 0 & 0 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \end{pmatrix}$$

Design an autoencoder with four hidden neurons to reconstruct the patterns, using gradient descent learning with a learning parameter $\alpha = 0.1$.

Find the weights, biases, hidden-layer activations and reconstructions of the input patterns at convergence.

Repeat the above by introducing a sparsity constraint with a penalty parameter $\beta = 0.5$ and sparsity parameter $\rho = 0.1$.

t10q1b.ipynb

Question 1

```
[7] # Display weights and biases
print(f'W:\n {autoencoder.W.data}\n')
print(f'W_prime:\n {autoencoder.W_prime.data}\n')
print(f'b:\n {autoencoder.b.data}\n')
print(f'b_prime:\n {autoencoder.b_prime.data}\n')
```

```
W:
tensor([[ 0.2693, -0.6963,  0.7908,  1.3101],
        [-0.8612, -0.7763,  2.1997, -0.2562],
        [-2.0359, -2.2188,  1.6613, -1.7386],
        [ 2.1135,  1.8982, -0.8923, -1.1621],
        [-2.6880, -2.7191, -2.8330,  0.7545],
        [ 0.0641,  0.3798,  0.0269, -0.2577],
        [ 1.0243,  0.4042, -0.9752, -3.5422],
        [-1.5969,  2.7576, -1.3955, -0.0790],
        [-1.9770,  1.5598, -2.8313,  2.1123]])
```

```
W_prime:
tensor([[ 2.7544, -1.6692, -3.3834,  4.2995, -3.1130, -0.8501,  2.5228, -0.7688,
         -2.0668],
        [ 3.2955, -2.3571, -4.0995,  5.2586, -4.9278, -1.1138,  3.3158,  5.5399,
         3.4433],
        [ 2.4002,  5.8303,  5.0095, -2.9082, -4.7423, -1.2832, -3.0701, -3.7878,
        -5.6506],
        [ 3.7825, -0.9833, -3.6115, -1.7625,  1.6002, -1.0698, -4.3225, -0.7160,
         4.1070]])
```

```
b:
tensor([-1.2380, -2.2536, -0.6383, -0.4749])
```

```
b_prime:
tensor([-1.0660, -2.9856,  1.0981, -2.2263,  1.9893, -4.4020,  1.6365, -3.0606,
        -1.3419])
```

```
# Evaluate result of reconstruction
with torch.no_grad():
    h, y, o = autoencoder(X)
print(f'Input:\n {X}\n')
print(f'Hidden activation:\n {h}\n')
print(f'Output:\n {y}\n')
print(f'Output_binary:\n {o}\n')
```

```
Input:
tensor([[1., 1., 1., 0., 0., 0., 0., 0., 0.],
        [1., 0., 0., 1., 0., 0., 1., 0., 0.],
        [0., 0., 1., 0., 1., 0., 1., 0., 0.],
        [1., 0., 0., 0., 1., 0., 0., 0., 1.],
        [1., 0., 0., 1., 0., 0., 1., 1., 1.]])
```

```
Hidden activation:
tensor([[0.0205, 0.0026, 0.9823, 0.2388],
        [0.8974, 0.3436, 0.1525, 0.0205],
        [0.0071, 0.0011, 0.0581, 0.0067],
        [0.0036, 0.0162, 0.0040, 0.9759],
        [0.1971, 0.9752, 0.0026, 0.1375]])
```

```
Output:
tensor([[9.0549e-01, 9.2173e-01, 9.9380e-01, 4.4885e-03, 8.5999e-02, 2.6298e-03,
         8.7004e-02, 9.5446e-04, 2.6120e-03],
        [9.5173e-01, 1.1840e-02, 6.5584e-02, 9.5070e-01, 3.9621e-02, 3.1252e-03,
         9.8883e-01, 8.0203e-02, 5.7781e-02],
        [2.9362e-01, 6.4900e-02, 7.9191e-01, 8.5439e-02, 8.4509e-01, 1.1085e-02,
         8.1012e-01, 3.6097e-02, 1.6060e-01],
        [9.3691e-01, 1.8599e-02, 7.6944e-02, 2.0678e-02, 9.6897e-01, 4.1843e-03,
         7.3666e-02, 2.4417e-02, 9.3653e-01],
        [9.6147e-01, 3.2267e-03, 1.7124e-02, 9.7066e-01, 3.8350e-02, 3.0001e-03,
         9.9154e-01, 8.8912e-01, 8.9647e-01]])
```

```
Output_binary:
tensor([[1., 1., 1., 0., 0., 0., 0., 0., 0.],
        [1., 0., 0., 1., 0., 0., 1., 0., 0.],
        [0., 0., 1., 0., 1., 0., 1., 0., 0.],
        [1., 0., 0., 0., 1., 0., 0., 0., 1.],
        [1., 0., 0., 1., 0., 0., 1., 1., 1.]])
```


Question 1

Comparing the sparsity:



Reconstruction loss only

```
Hidden activation:
tensor([[0.7265, 0.0868, 0.9958, 0.0081],
        [0.9890, 0.1727, 0.3677, 0.9023],
        [0.0207, 0.0045, 0.9912, 0.9852],
        [0.0226, 0.9916, 0.9659, 0.7991],
        [0.9845, 0.9704, 0.0123, 0.9855]])
```

Reconstruction loss + Sparsity loss

```
Hidden activation:
tensor([[0.0205, 0.0026, 0.9823, 0.2388],
        [0.8974, 0.3436, 0.1525, 0.0205],
        [0.0071, 0.0011, 0.0581, 0.0067],
        [0.0036, 0.0162, 0.0040, 0.9759],
        [0.1971, 0.9752, 0.0026, 0.1375]])
```


Question 1



Question 2


Create 100 images of 10x10 size by randomly generating pixel values between 0.0 and 1.0 from a uniform distribution.

Design the following autoencoders to reconstruct the input patterns, using mean square error as the cost function:

- a. An undercomplete autoencoder with 49 hidden neurons
- b. An overcomplete autoencoder with 144 hidden neurons
- c. A sparse autoencoder with 144 hidden neurons and training with sparsity parameter $\rho = 0.05$ and penalty parameter $\beta = 0.5$.

Compare features learned by different autoencoders.

Question 2a




Question 2a




10x10 weights of 49 filters

Question 2a




Input



Reconstruction

Question 2a



7x7 hidden activations (reshaped from 49 dimensions) of 100 samples

Question 2

Create 100 images of 10x10 size by randomly generating pixel values between 0.0 and 1.0 from a uniform distribution.


Design the following autoencoders to reconstruct the input patterns, using mean square error as the cost function:

- a. An undercomplete autoencoder with 49 hidden neurons
- b. An overcomplete autoencoder with 144 hidden neurons
- c. A sparse autoencoder with 144 hidden neurons and training with sparsity parameter $\rho = 0.05$ and penalty parameter $\beta = 0.5$.


Compare features learned by different autoencoders.

t10q2b.ipynb

Question 2b




Question 2b




10x10 weights of 144 filters

Question 2b




Input



Reconstruction

Question 2b



12x12 hidden activations (reshaped from 144 dimensions) of 100 samples

Question 2


Create 100 images of 10x10 size by randomly generating pixel values between 0.0 and 1.0 from a uniform distribution.

Design the following autoencoders to reconstruct the input patterns, using mean square error as the cost function:


- a. An undercomplete autoencoder with 49 hidden neurons
- b. An overcomplete autoencoder with 144 hidden neurons
- c. A sparse autoencoder with 144 hidden neurons and training with sparsity parameter $\rho = 0.05$ and penalty parameter $\beta = 0.5$.

Compare features learned by different autoencoders.

Question 2c




Question 2c




10x10 weights of 144 filters

Question 2c




Input



Reconstruction


Question 2c



12x12 hidden activations (reshaped from 144 dimensions) of 100 samples


Question 2c

Without sparsity



12x12 hidden activations (reshaped from 144 dimensions) of 100 samples

With sparsity



12x12 hidden activations (reshaped from 144 dimensions) of 100 samples

Question 3

Design a denoising autoencoder to reconstruct MNIST images:

<http://yann.lecun.com/exdb/mnist/>

- (a) Assume one hidden layer with 625 neurons, multiplicative noise, and cross-entropy cost function. Use 10% corruption level, learning factor $\alpha = 0.1$, batch size = 128, sparsity constant $\rho = 0.02$, and penalty parameter $\beta = 0.4$.

Plot the learning curves, the weights, and the hidden layer activations for sample test images.

- (b) Add another hidden layer with 100 neurons and train the autoencoder as before. Plot the feature maps.

Plot the learning curves, the weights, and the hidden layer activations for sample test images

- (c) Add a softmax layer on top of the second hidden layer to design a classifier. Show learning curves and find the accuracy of the classifier.


Plot the learning curves and the weights and find the accuracy for test patterns.

t10q3.ipynb


Deep stacked autoencoders



Training first
hidden layer




Training second
hidden layer.
Reconstructs
first-hidden
layer output




Autoencoder
with two
hidden layers

Deep stacked autoencoders



Training the third hidden layer.
Reconstructs second hidden
layer output




Autoencoder with
three hidden layers

Semi-Supervised Classification

- Many images, but few ground truth labels

start unsupervised
train autoencoder on many images



supervised fine-tuning
train classification network on labeled images

